

What is “Optimisation”?

When is it “Premature”?

Version 1.1-beta.1

Dr. Colin Hirsch

What is “Premature Optimisation”?

Why is it the “Root of all Evil”?

Version 1.1-beta.1

Dr. Colin Hirsch

"We should forget about small efficiencies, say about 97% of the time: premature optimization is the root of all evil. Yet we should not pass up our opportunities in that critical 3%"

Donald Knuth

“

all evil

premature optimization is the root of

”

Donald Knuth

Optimising what?

- Time vs. space?
 - Latency vs. throughput?
 - Average vs. worst case?
 - Performance vs. efficiency?
 - One vs. few vs. many threads?
-
- Development time?
 - Number of unit tests?
 - Maintainability of code?

What is “optimisation”?

1st attempt

Making code faster!

Why do we optimise?

- Efficient code is **green** code
 - saves energy, money and resources
- Fast code makes for **better user experience**
 - higher user satisfaction, more sales
- Some code has **performance requirements**
 - embedded and real-time systems, ...
- It's **fun** and **gratifying** to make code faster
 - which is why we *might* have a tendency to optimise prematurely

So where's the “evil”?

“Evil” optimisation side-effects

- Adding code and complexity
- Making code more error prone
- Getting lost in micro optimisations
- Making things slower instead of faster
- Breaking the functionality in corner cases
- Making the code less flexible and malleable
- Wasting time optimising in wrong places
- Making the code harder to read and understand
- Requiring more (unit) tests to verify correctness

What is the “evil”?

- Adding code and complexity
- Managing complexity is *a key challenge* of software development
 - Layered models
 - Divide and conquer
 - Functions and classes
 - Modules and interfaces
 - Other abstraction mechanisms
 - *Less is more* (YAGNI, DRY, minimalistic code)

Big picture: taoCONFIG vs. PEGTL

- taoCONFIG reads configuration files for applications
 - Configuration files are small and read once
 - Optimising taoCONFIG is not our priority
- PEGTL parses data according to a user's grammar
 - Some users parse a lot of data and/or a lot of times
 - PEGTL performance is always on our minds
 - But we are not implementing packrat parsing!
- Takeaway: Look at the big picture before optimising!

Hearsay based optimisation: $O(n)$

- PEGTL recursive-descent approach is $O(n^2)$ or worse
- “Packrat parsing (with memoisation) is sooo much better with its $O(n)$ ”
- Yes, packrat parsing has better worst-case complexity, but:
 - A much higher constant factor (overhead), and:
 - has more code and uses more memory at runtime, and:
 - how many real-world grammars hit the worst case anyhow?
- Nobody has convinced us yet that the PEGTL would benefit from packrat
- Takeaway: Theoretical advantages don't always hold in practice!

Hearsay based optimisation: virtual

- “Virtual functions are slow”
- Yes, they are slower than plain functions, but:
 - Are they slower than the alternatives?
 - How complicated are these alternatives?
- Runtime polymorphism has *some* cost!
- Virtual functions should be used where *appropriate*,
 - and [their overhead] avoided where not necessary
- Takeaway: Compare and choose wisely!

Statistics guided optimisation

- Needed to extend app to keep track of certain things
- Question was which container to use for these things
- Use cases were discussed and access patterns analysed
- Then we discovered the number of these things at any given time
 - Nearly always either 0 and 1
- In other words it doesn't matter which data structure is used
- Takeaway: Know what you are optimising for!

Library vs. hand rolled: taoJSON

- taoJSON value class holds different types
- Initially based on union & enum & switch statements (**fast!**)
- Later changed to std::variant (**slow?**)
- That's the opposite of optimising, but:
 - Pages of low-level code were removed (**great!**)
 - Performance did not suffer noticeably (**good!**)
- Takeaway: The standard library is often very good and/or good enough!

Hand optimised: FLC video player

- Once upon a time I had an FLC video file on my Amiga
- Found a player written in *assembly*
 - Hand-optimised read-and-decode loop (**fast but ... assembly!**)
- But the host adapter can DMA from HDD to RAM...
- Wrote a multi-threaded player in C
 - Use CPU to decode during asynchronous DMA (**faster and ... easier!**)
- Takeaway: Restructuring on high-level beats low-level optimisations!

Just doing our jobs: “Good code”

- Writing *appropriate / elegant / minimalistic* code
 - `std::unique_ptr` vs. `std::shared_ptr`
 - `std::vector` vs. `std::list` vs. `std::deque` vs. `std::set`
 - passing by value or by reference
- Reasoning about these choices is reasoning about structure and design!
- And these choices convey information to the reader!
- Takeaway: Not everything that *optimises* is an *optimisation*!

Just doing our jobs!

- Everything that simplifies code or reduces complexity
 - Even if it makes the code faster as side effect
- Everything that makes code more readable and maintainable
 - Even if it makes the code faster as side effect
- Making the structure of the code match the structure of the problem
 - Frequently produces *good* or at least *good enough* performance
- Most things that make code faster without increasing complexity
 - Choosing the most efficient alternative without drawbacks

Complexity vs. optimisation?

What is “optimisation”?

2nd attempt

Making code faster...

...while increasing complexity!

What about “premature”?

Premature optimisation checklist

1st attempt

- Am I optimising the right places?
 - Probably not, the profiler is your friend!
 - Is it even worth it, is the code run often enough?
- Will my change improve performance?
 - Benchmark a prototype or mockup or something!
- Am I optimising code that will survive?
 - Is the feature needed in the first place?
 - Will a higher-level improvements eliminate the code?
- Will my unit tests catch bugs introduced while optimising?
- Am I keeping in mind that the most efficient code is ... no code?

Recommended approach

- First create the baseline
 - **Correct** code
 - Nice, simple, minimalistic, elegant, ... code
 - This is usually quite fast/efficient,
 - perhaps even fast/efficient enough
- Then think about what might need optimising
 - And use the profiler and the questions to be sure
 - Low hanging fruit with local impact can be fair game

What is “premature optimisation”?

Optimising code...

...before the baseline!

baseline = correct, clean, elegant, minimalistic code

...before doing the checklist!

checklist = all the questions from two slides ago

Thank you!

Colin Hirsch — mail@cohi.at